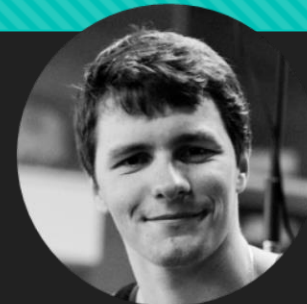




**How about no grep and zabbix?**

**or**

**ELK based metrics and alerts.**



Vladimir Pavkin  
QIWI Conf #3



Scala Developer



gettopical.com

# Plan

- ELK stack components
- ELK setups
- Basic usage (grep)
- Advanced usage (alerts, metrics)

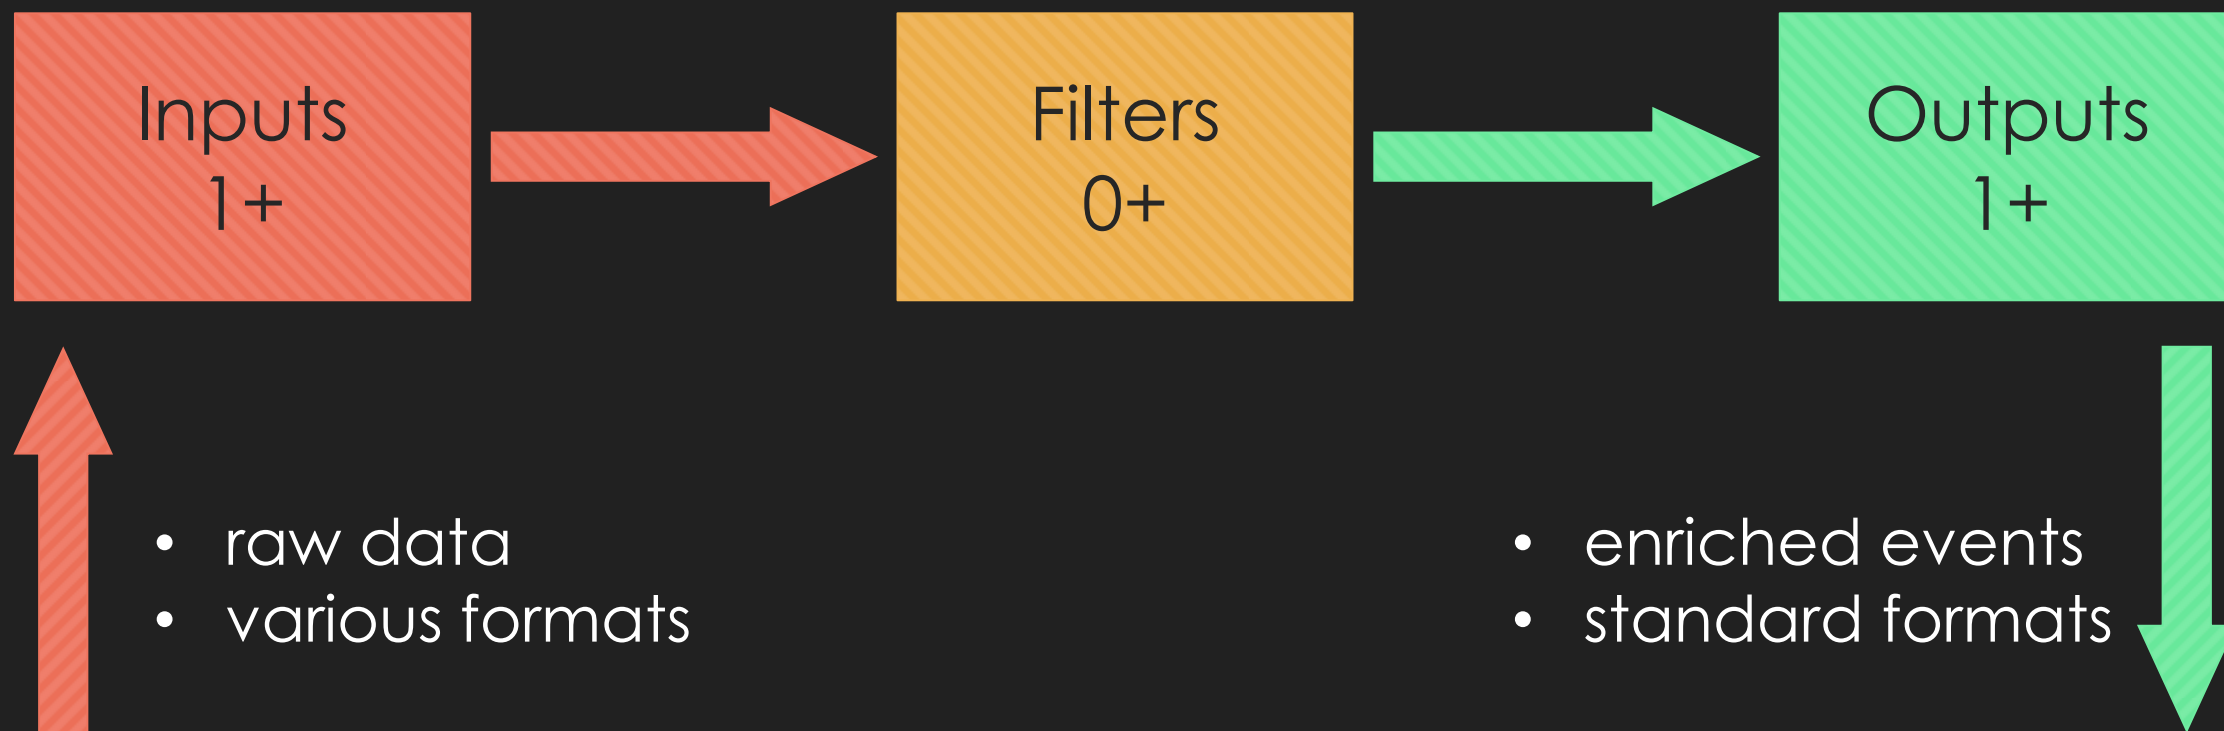
# Da heck is ELK?

Elasticsearch

Logstash

Kibana

# Logstash



# Logstash inputs

- file, syslog, log4j
- http, tcp, udp
- rabbitmq
- twitter, rss
- ...

# Logstash filters

- grok
- geoip
- drop, clone, mutate, split
- ...

# Logstash outputs

- stdout, file
- elasticsearch
- rabbitmq
- http, tcp, udp
- ...



# Logstash transformation example

```
127.0.0.1 - - [11/Dec/2013:00:01:45 -0800]  
"GET /xampp/status.php HTTP/1.1" 200 3891  
"http://cadenza/xampp/navi.php" "Mozilla/5.0 (Macintosh;  
Intel Mac OS X 10.9; rv:25.0) Gecko/20100101 Firefox/25.0"
```

# Logstash transformation example

```
{  
  "message" : "127.0.0.1 - - [11/Dec/2013:00:01:45 -0800] \"GET /xampp/status.php  
HTTP/1.1\" 200 3891 \"http://cadenza/xampp/navi.php\" \"Mozilla/5.0 (Macintosh;  
Intel Mac OS X 10.9; rv:25.0) Gecko/20100101 Firefox/25.0\\\"\",  
  "@timestamp" : "2013-12-11T08:01:45.000Z",  
  "@version" : "1",  
  "host" : "cadenza",  
  "clientip" : "127.0.0.1",  
  "path" : "/xampp/status.php",  
  "response" : "200",  
  // ...  
}
```

# Elasticsearch



Distributed data storage, optimized for text search.

Is managed/queried through HTTP REST API.

Based on open-source Apache Lucene project.

# Elasticsearch vocab

- **Document** - data unit used for indexing and searching. Contains **fields**. (JSON)
- **Field** - a part of a **Document**, has a **name** and a **value**.
- **Term** - the unit of search. A word we look for in text

# Elasticsearch inverted index

## Documents:

1 => How about no grep

2 => How about no zabbix

## Index:

how => <1>, <2>

about => <1>, <2>

zabbix => <2>

...

# Elasticsearch query DSL

- `message:exception`
- `log_level:(ERROR OR INFO)`
- `stack_trace:"Out of memory"`
- `_exists_:body`
- `date:[now-1h TO now]`
- `age:(>=10 AND <20)`
- wildcards (\*, ?)
- regex search
- fuzzy search (Levenshtein)
- proximity search
- boosting weights

# Elasticsearch query DSL

Even more power with JSON query DSL:

```
{
  "query_string" : {
    "fields" : ["message", "stack_trace"],
    "query" : "exception AND fatal"
  }
}
```

# Elasticsearch clustering

Just easy as hell





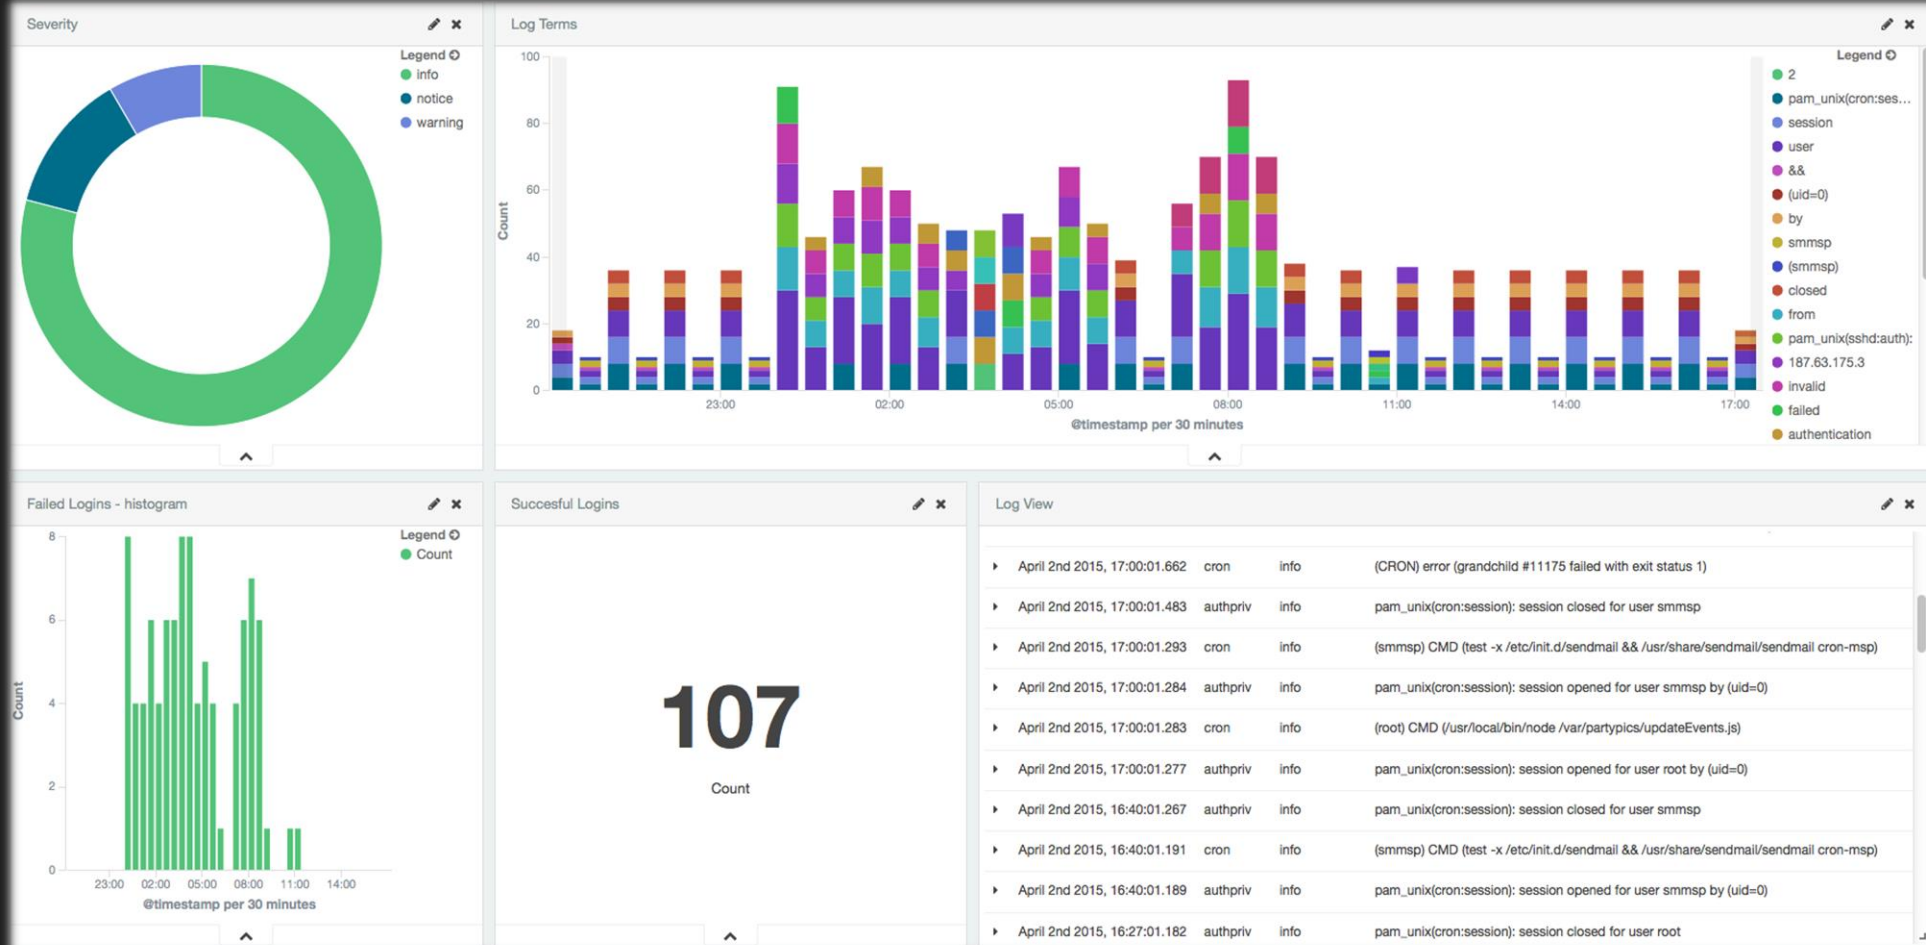
Converts these:

```
{  
  "message" : "127.0.0.1 - - [11/Dec/2013:00:01:45 -0800] \"GET /xampp/status.php  
HTTP/1.1\" 200 3891 \"http://cadenza/xampp/navi.php\" \"Mozilla/5.0 (Macintosh; Intel  
Mac OS X 10.9; rv:25.0) Gecko/20100101 Firefox/25.0\"",  
  "@timestamp" : "2013-12-11T08:01:45.000Z",  
  "@version" : "1",  
  "host" : "cadenza",  
  "clientip" : "127.0.0.1",  
  "path" : "/xampp/status.php",  
  "response" : "200",  
  // ...  
}
```

# Kibana



To this:



# Kibana timestamp field

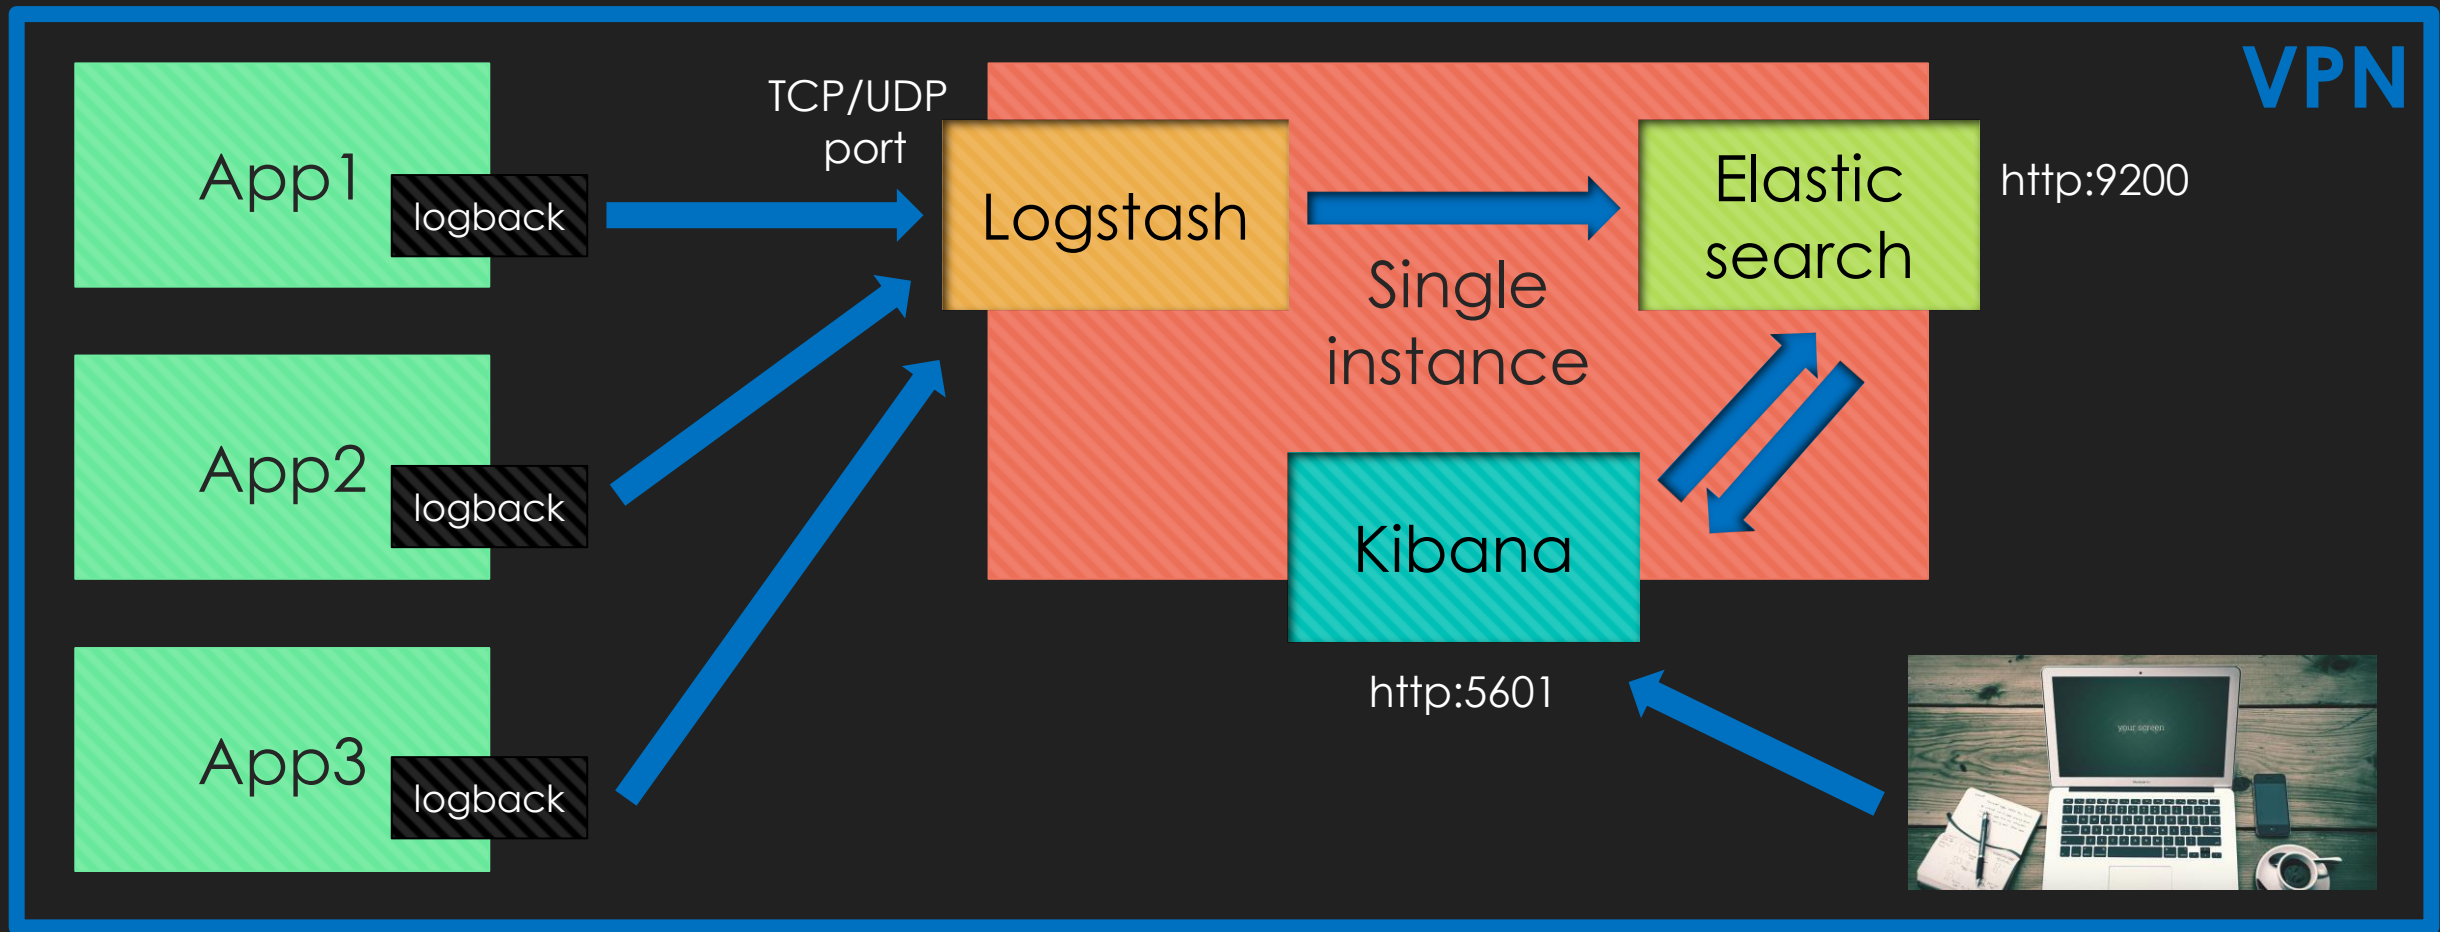
Not strings/objects

Events!

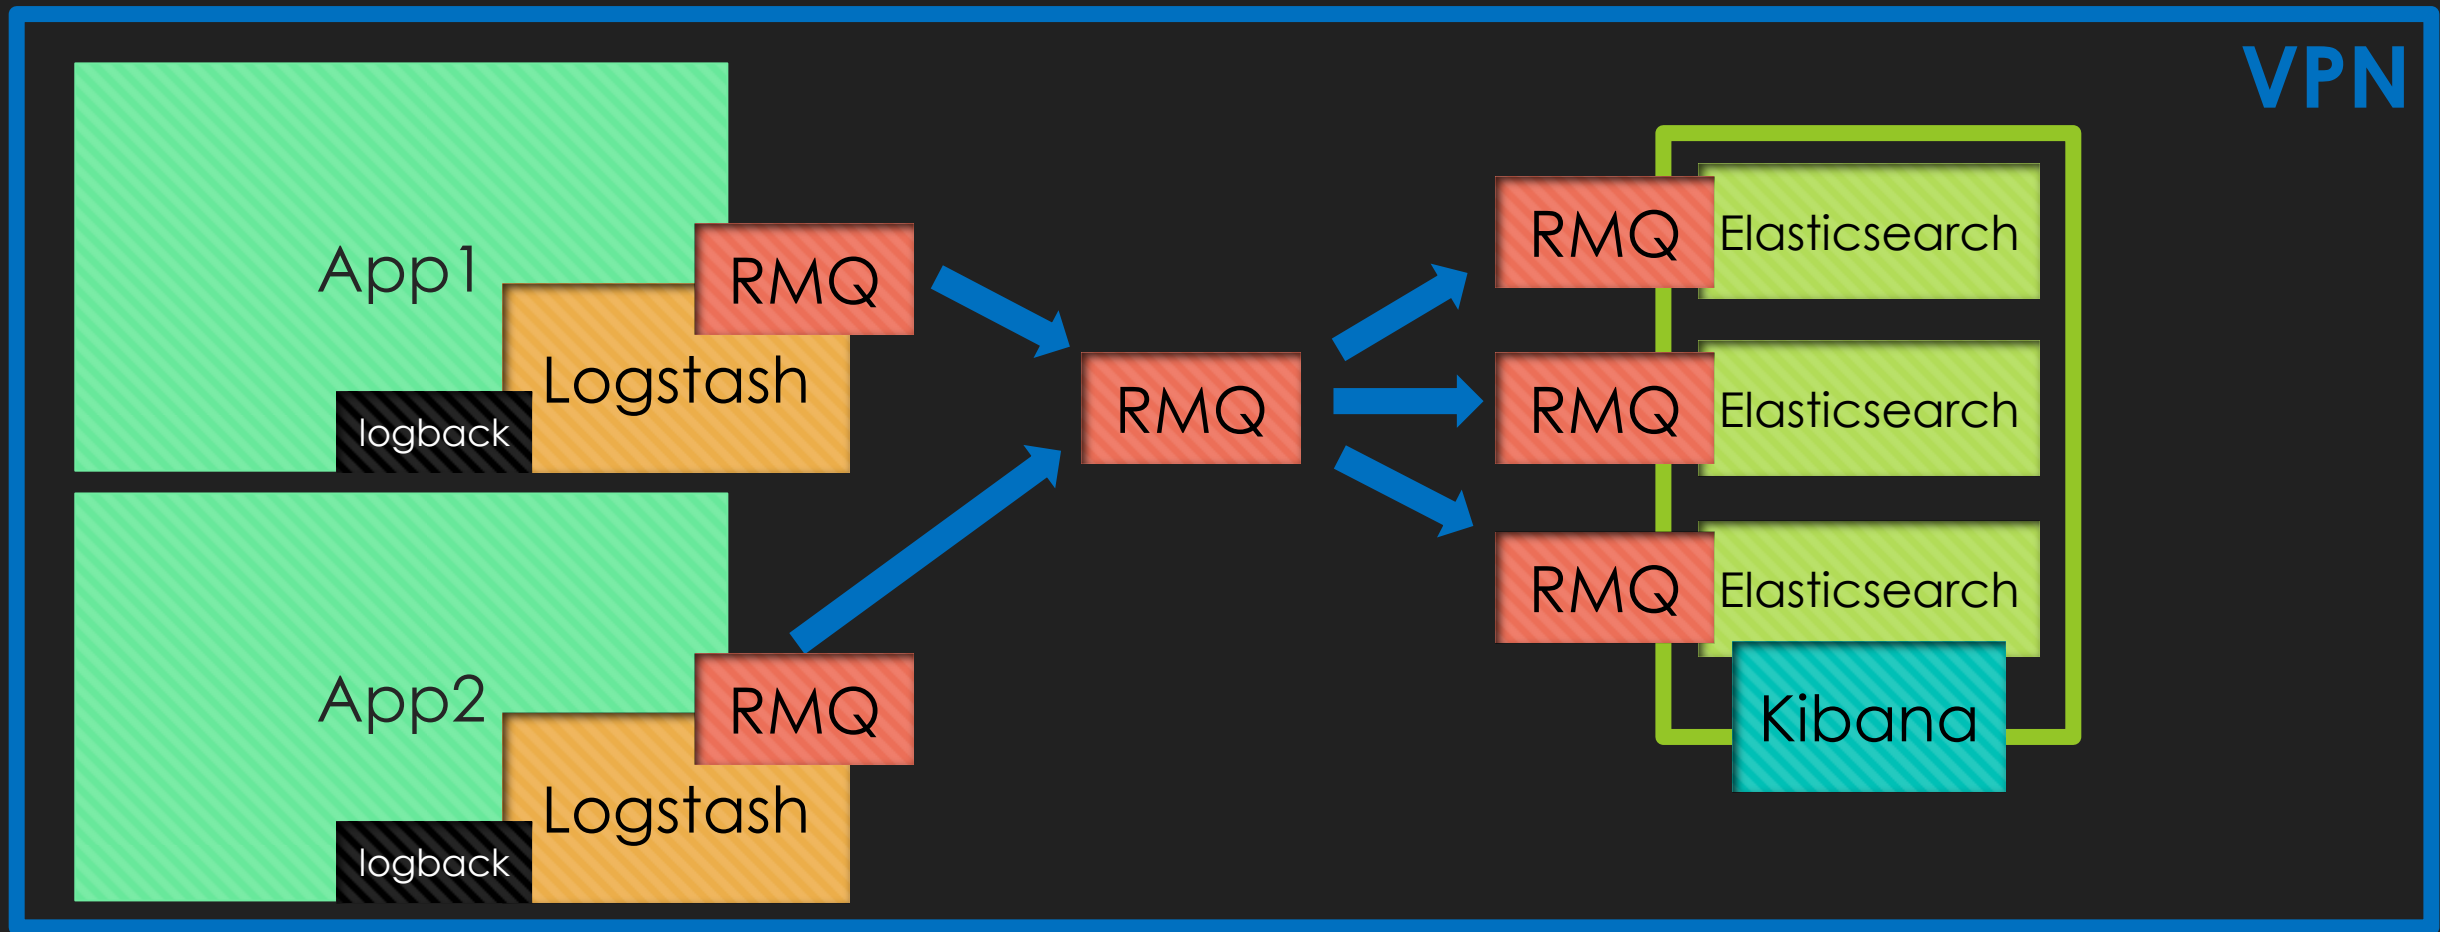


# ELK Setups

# Simple setup (startup mode)



# Complex setup (QIWI)



# Basic ELK usage scenarios



# 1. Search that whole cluster!

- Events from all your apps are in one place!
- Rocking useful and flexible timeline
- Powerful, text oriented search engine



## 2. Default logback fields

- Quickly filter events by:
  - Host name or IP
  - Log level, logger name
- Search only in stack trace.

# 3. Customize and save searches

- Save frequently used search
- Show only specific fields  
(e.g. message and host name)

Showtime

# Advanced ELK. Metrics and alerts.

# 1. Custom fields (search/filter)

Payment event:

```
{  
  "user"      : "9112223344",  
  "amount"    : 500.00,  
  "provider"  : 464  
}
```

# Log custom fields (JVM)

```
libraryDependencies +=  
  "net.logstash.logback" % "logstash-logback-encoder" % "4.5.1"  
  
<appender name="LOGSTASH"  
class="net.logstash.logback.appender.LogstashTcpSocketAppender">  
  <destination>127.0.0.1:4560</destination>  
</appender>  
  
<root level="DEBUG">  
  <!--...-->  
  <appender-ref ref="LOGSTASH" />  
</root>
```

# Log custom fields (JVM)

```
import net.logstash.logback.marker.Markers._

val fields = Map[String, Any](
  "str_field" -> "str_value",
  "num_field" -> 666
)

logger.info(appendEntries(fields), "log message");
```

## 2. Elastalert



- Configured with a set of “rules”
- Performs periodical queries to ES
- Triggers alerts for matched rules



## 2. Elastalert rule types

- Frequency
- Spike
- Flatline
- Cardinality
- Blacklist, Whitelist
- Arbitrary query

```
# rules/500_frequency.yaml
es_host: localhost
es_port: 9200
name: 500 code frequency alert
type: frequency
index: logstash-*
num_events: 50
timeframe:
  minutes: 1
alert:
  - "email"
email:
  - "s.solonin@qiwi.ru"
```

## 2. Elastalert alert types

- Email
- Slack
- Jira
- Custom shell command

## 2. Elastalert with custom fields

Any alert you can imagine can be implemented on custom fields:

- Average payment drops below 300 RUB
- Processing response latency spike
- User A buys something on Herbalife

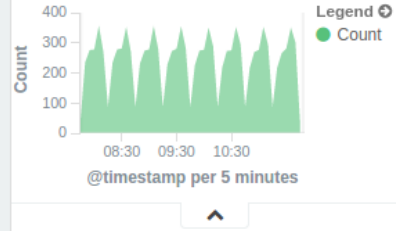
# 3. Monitoring and Metrics

- Realtime
- Lots of aggregations
- Beautiful visualizations
- Super power with custom fields

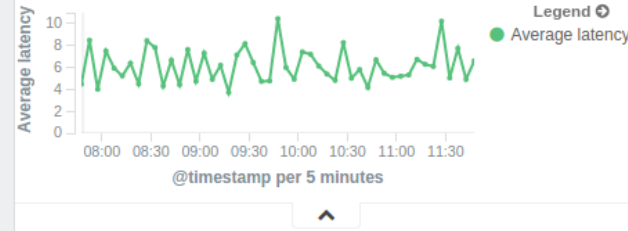
Fetched recent tweets [Graph]



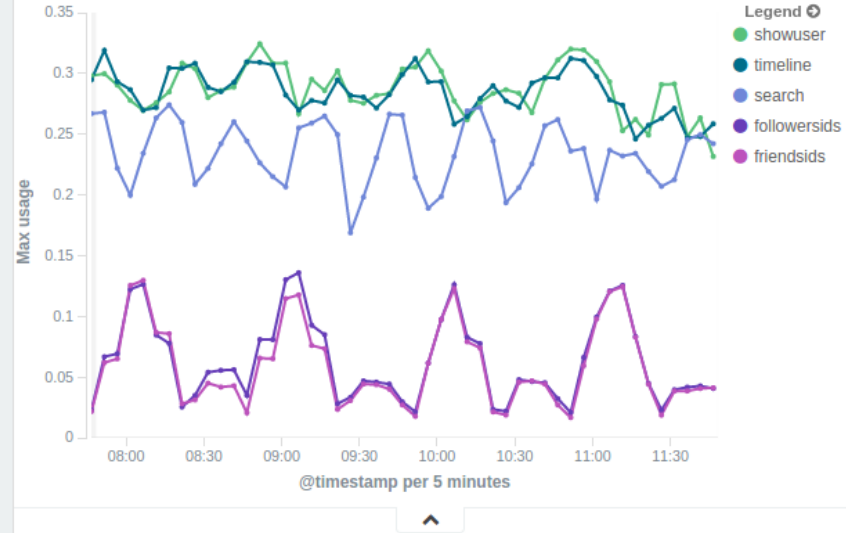
New ExpandedURL saved [Graph]



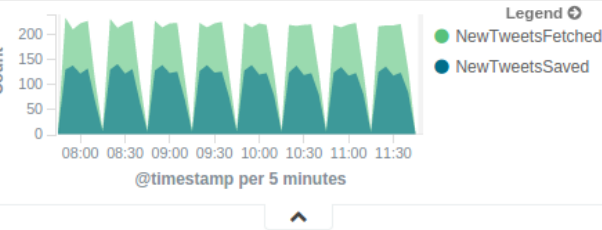
Twitter Utility Account Manager Actor Latency



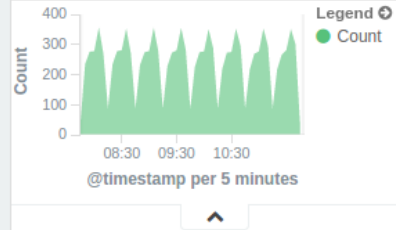
Utility Pool Exhauster



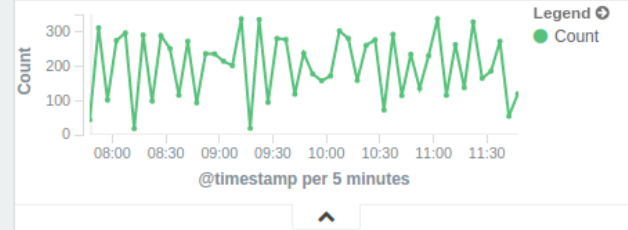
Fetched popular tweets [Graph]



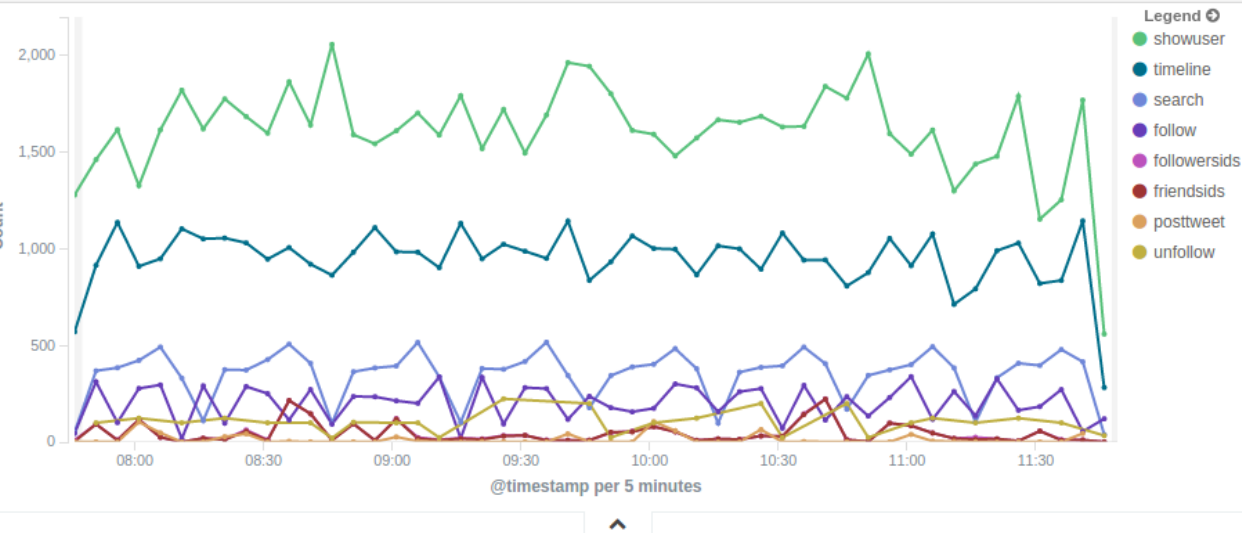
New ShortURL saved [Graph]



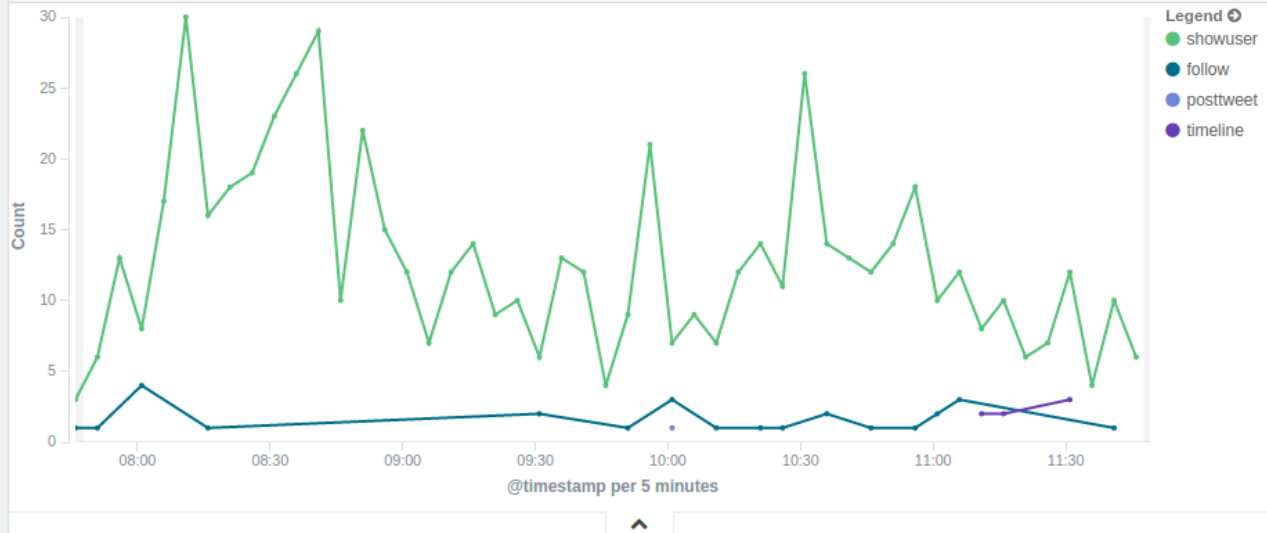
Number of follows (total)

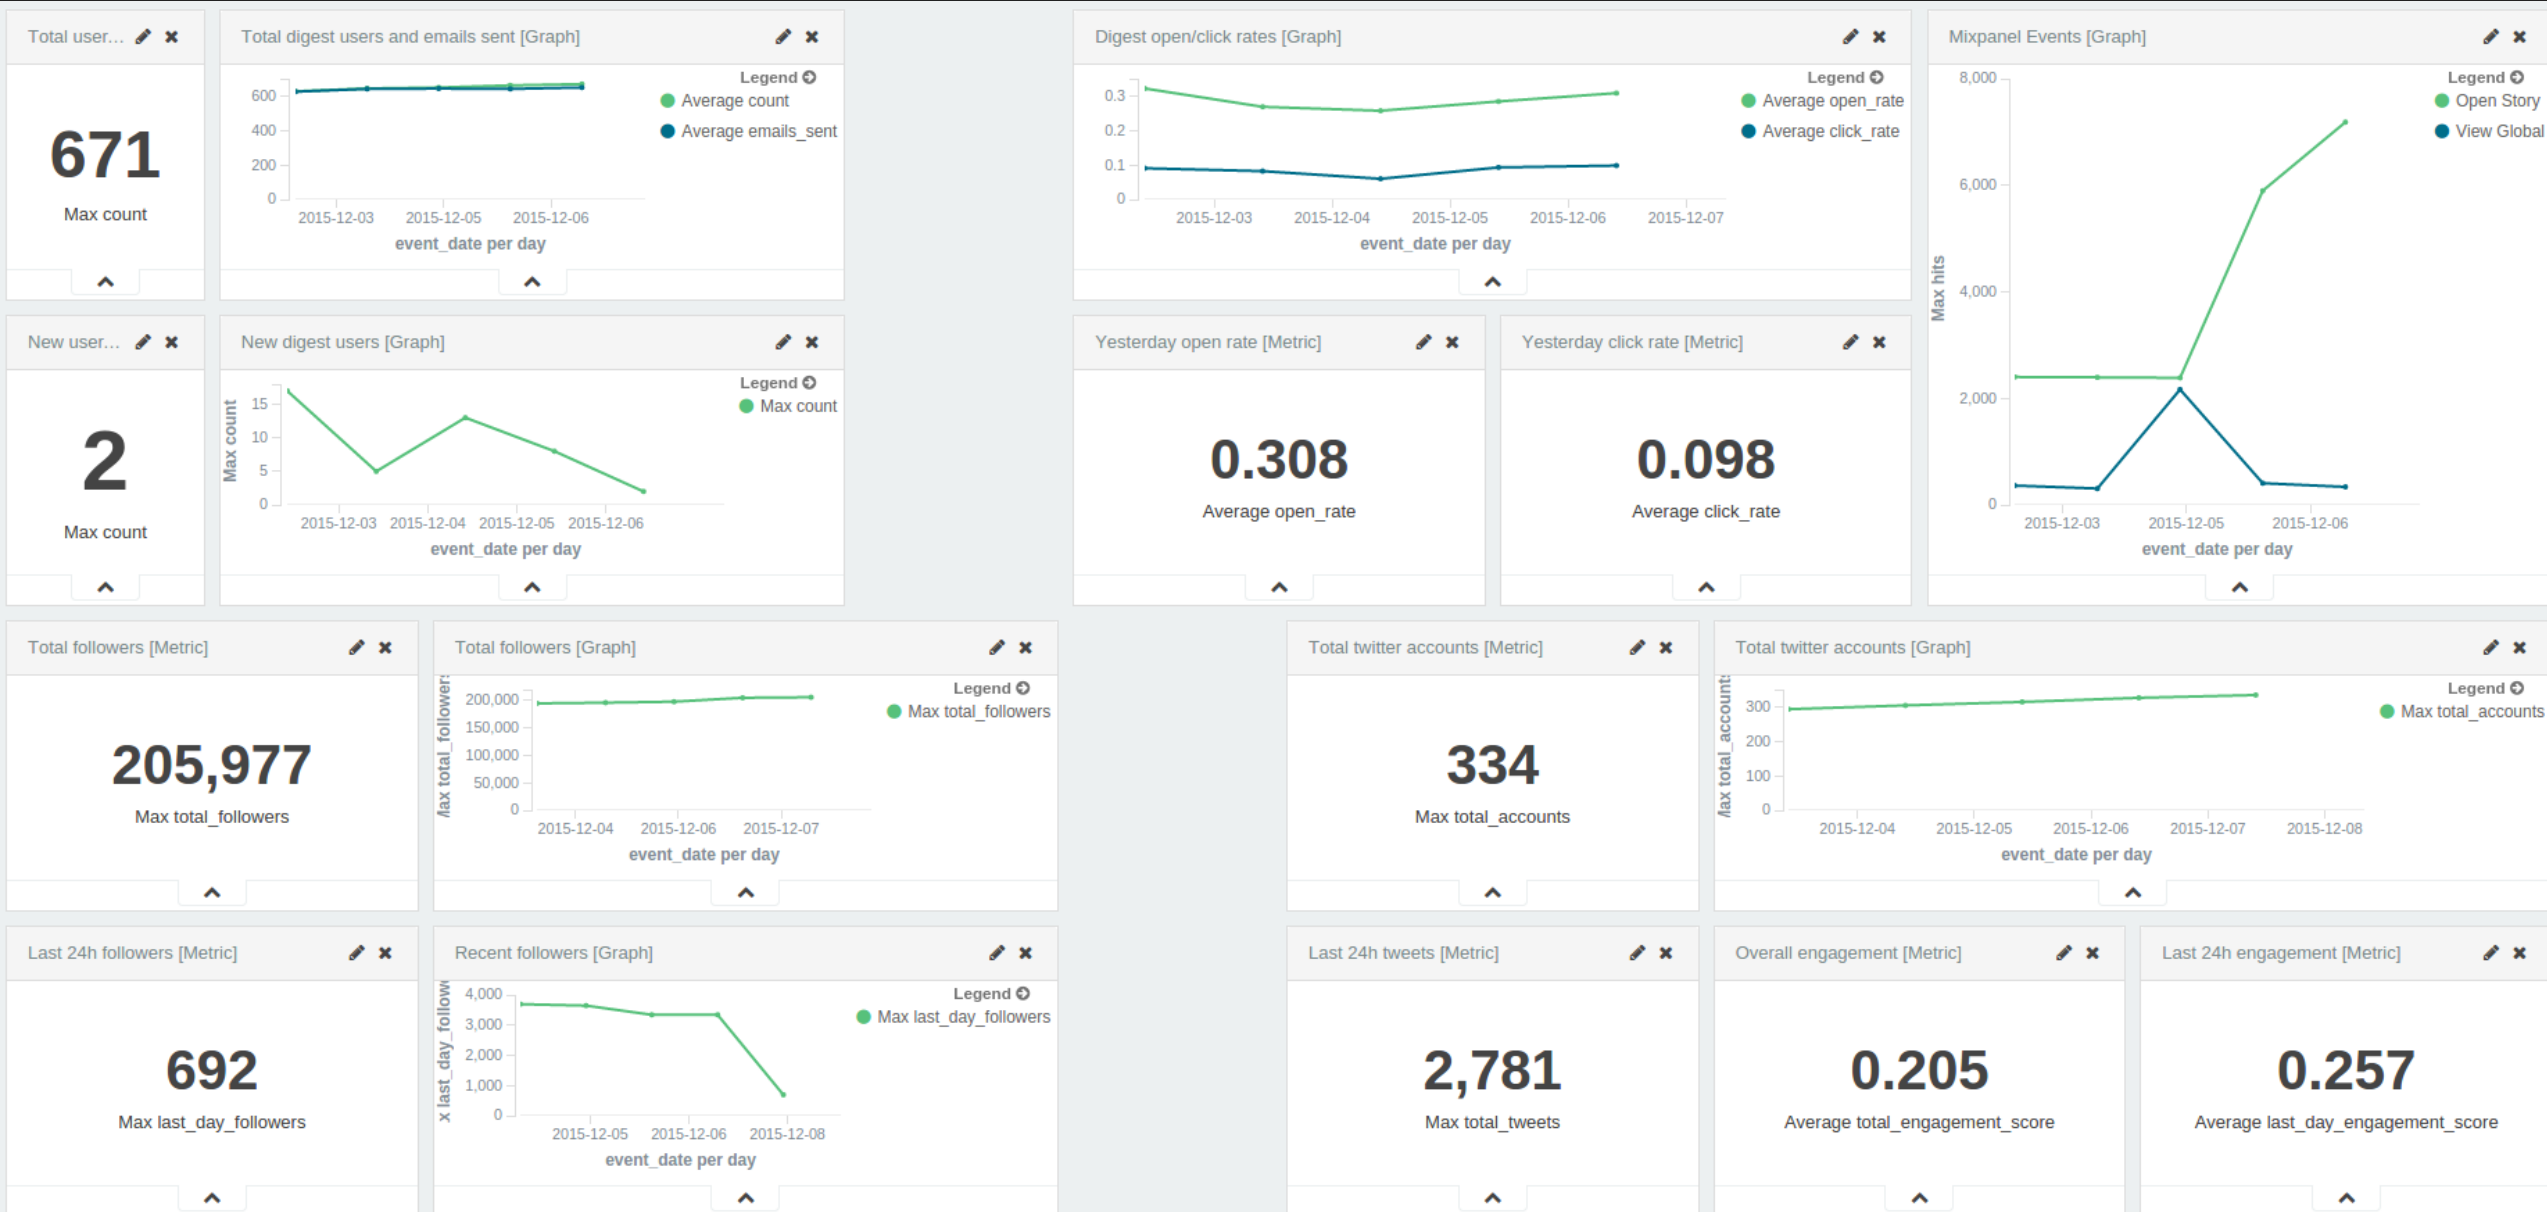


Twitter API requests



Twitter API errors





Showtime

### 3. Why so useful?



- Tech monitoring (UI and flexibility)
- Business metrics (approx.)
- Alerts (both tech and business)
- Instant market/system reaction



# Thanks !

